

assignment-5

February 19, 2024

```
[21]: import pandas as pd
import numpy as np
```

```
[22]: df = pd.read_csv("/content/Social_Network_Ads.csv")
```

```
[23]: df.head()
```

```
[23]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19.0	19000.0	0
1	15810944	Male	35.0	20000.0	0
2	15668575	Female	26.0	43000.0	0
3	15603246	Female	27.0	57000.0	0
4	15804002	Male	19.0	76000.0	0

```
[24]: df.info()
```

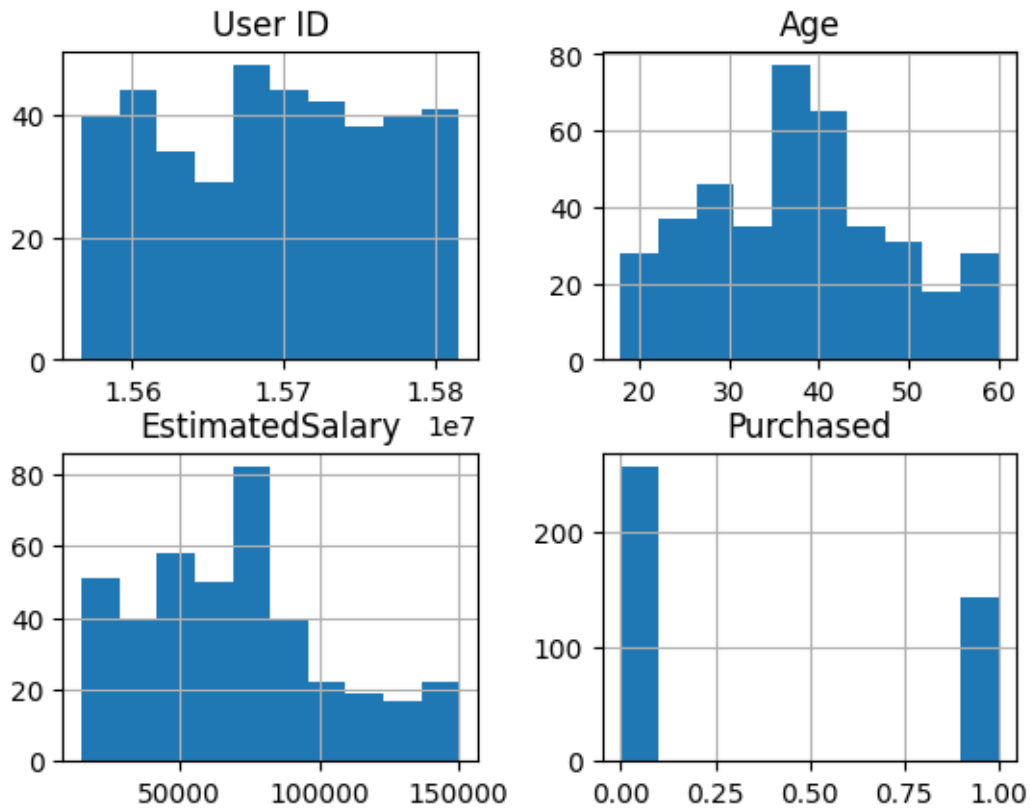
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   User ID                400 non-null   int64  
1   Gender                 400 non-null   object  
2   Age                   400 non-null   float64 
3   EstimatedSalary        400 non-null   float64 
4   Purchased              400 non-null   int64  
dtypes: float64(2), int64(2), object(1)
memory usage: 15.8+ KB
```

```
[25]: df.isnull().sum()
```

```
[25]: User ID                0
Gender                  0
Age                    0
EstimatedSalary        0
Purchased              0
dtype: int64
```


```
[26]: df.hist()
```

```
[26]: array([[<Axes: title={'center': 'User ID'}>,  
          <Axes: title={'center': 'Age'}>],  
          [<Axes: title={'center': 'EstimatedSalary'}>,  
          <Axes: title={'center': 'Purchased'}>]], dtype=object)
```




```
[27]: df['Age'].plot.box()
```

```
[27]: <Axes: >
```




```
[28]: df['EstimatedSalary'].plot.box()
```

```
[28]: <Axes: >
```



```
[29]: df['Age'].hist()
```

```
[29]: <Axes: >
```



```
[30]: X = df.drop('Purchased',axis=1)
```

```
[31]: X
```

```
[31]:
```

	User ID	Gender	Age	EstimatedSalary
0	15624510	Male	19.0	19000.0
1	15810944	Male	35.0	20000.0
2	15668575	Female	26.0	43000.0
3	15603246	Female	27.0	57000.0
4	15804002	Male	19.0	76000.0
..
395	15691863	Female	46.0	41000.0
396	15706071	Male	51.0	23000.0
397	15654296	Female	50.0	20000.0
398	15755018	Male	36.0	33000.0
399	15594041	Female	49.0	36000.0

```
[400 rows x 4 columns]
```

```
[32]: Y = df['Purchased']
```

```
[33]: Y
```

```
[33]: 0      0
      1      0
      2      0
      3      0
      4      0
      ..
      395    1
      396    1
      397    1
      398    0
      399    1
      Name: Purchased, Length: 400, dtype: int64
```

```
[34]: X.drop('Gender',inplace = True,axis=1)
```

```
[35]: X
```

```
[35]:      User ID   Age  EstimatedSalary
0    15624510  19.0         19000.0
1    15810944  35.0         20000.0
2    15668575  26.0         43000.0
3    15603246  27.0         57000.0
4    15804002  19.0         76000.0
..         ...   ...
395  15691863  46.0         41000.0
396  15706071  51.0         23000.0
397  15654296  50.0         20000.0
398  15755018  36.0         33000.0
399  15594041  49.0         36000.0
```

```
[400 rows x 3 columns]
```

```
[36]: X.drop('User ID',inplace = True,axis=1)
```

```
[37]: X
```

```
[37]:      Age  EstimatedSalary
0    19.0         19000.0
1    35.0         20000.0
2    26.0         43000.0
3    27.0         57000.0
4    19.0         76000.0
..     ...
395  46.0         41000.0
396  51.0         23000.0
397  50.0         20000.0
398  36.0         33000.0
```

```
399 49.0          36000.0
```

```
[400 rows x 2 columns]
```

```
[38]: from sklearn.model_selection import train_test_split
```

```
[39]: x_train,x_test,y_train,y_test = train_test_split(X,Y,test_size=0.  
↳25,random_state=0)
```

```
[40]: x_test
```

```
[40]:      Age  EstimatedSalary  
132  30.0          87000.0  
309  38.0          50000.0  
341  35.0          75000.0  
196  30.0          79000.0  
246  35.0          50000.0  
..    ...           ...  
146  27.0          96000.0  
135  23.0          63000.0  
390  48.0          33000.0  
264  48.0          90000.0  
364  42.0         104000.0
```

```
[100 rows x 2 columns]
```

```
[41]: x_train
```

```
[41]:      Age  EstimatedSalary  
250  44.0          39000.0  
63   32.0         120000.0  
312  38.0          50000.0  
159  32.0         135000.0  
283  52.0          21000.0  
..    ...           ...  
323  48.0          30000.0  
192  29.0          43000.0  
117  36.0          52000.0  
47   27.0          54000.0  
172  26.0         118000.0
```

```
[300 rows x 2 columns]
```

```
[42]: y_test
```

```
[42]: 132    0  
309    0
```

```

341    0
196    0
246    0
..
146    1
135    0
390    1
264    1
364    1
Name: Purchased, Length: 100, dtype: int64

```

```
[43]: y_test
```

```

[43]: 132    0
      309    0
      341    0
      196    0
      246    0
      ..
      146    1
      135    0
      390    1
      264    1
      364    1
Name: Purchased, Length: 100, dtype: int64

```

```
[44]: y_train
```

```

[44]: 250    0
      63    1
      312    0
      159    1
      283    1
      ..
      323    1
      192    0
      117    0
      47     0
      172    0
Name: Purchased, Length: 300, dtype: int64

```

```
[45]: from sklearn.preprocessing import StandardScaler
      std = StandardScaler()
```

```
[46]: x_train = std.fit_transform(x_train)
      x_test = std.fit_transform(x_test)
```



```
[47]: from sklearn.linear_model import LogisticRegression
```

```
[48]: model = LogisticRegression()
```

```
[49]: model.fit(x_train,y_train)
```

```
[49]: LogisticRegression()
```

```
[50]: y_pred = model.predict(x_test)
```

```
[51]: from sklearn.metrics import   
      ↪confusion_matrix,accuracy_score,precision_score,recall_score  
CM = confusion_matrix(y_test,y_pred)  
print(accuracy_score(y_test,y_pred))
```

0.87

```
[52]: CM
```

```
[52]: array([[63,  5],  
          [ 8, 24]])
```

```
[52]:
```

```
[53]: print(precision_score(y_test,y_pred))
```

0.8275862068965517

```
[54]: print(recall_score(y_test,y_pred))
```

0.75

```
[55]: from sklearn.metrics import classification_report  
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.89	0.93	0.91	68
1	0.83	0.75	0.79	32
accuracy			0.87	100
macro avg	0.86	0.84	0.85	100
weighted avg	0.87	0.87	0.87	100

```
[56]: myData = [[45,85000],[35,65000],[35,120000]]
```

```
[57]: myData
```

```
[57]: [[45, 85000], [35, 65000], [35, 120000]]
```

```
[58]: myData = std.fit_transform(myData)
```

```
[59]: myData
```

```
[59]: array([[ 1.41421356, -0.21997067],  
        [-0.70710678, -1.09985336],  
        [-0.70710678,  1.31982404]])
```

```
[60]: dataPred = model.predict(myData)
```

```
[61]: dataPred
```

```
[61]: array([1, 0, 0])
```

```
[62]: #using min max scaling  
from sklearn.preprocessing import MinMaxScaler  
min_max_scale = MinMaxScaler()
```

```
[63]: x_train = min_max_scale.fit_transform(x_train)  
x_test = min_max_scale.fit_transform(x_test)
```

```
[64]: model1= LogisticRegression()
```

```
[65]: model1.fit(x_train,y_train)
```

```
[65]: LogisticRegression()
```

```
[66]: y_pred = model.predict(x_test)
```

```
[67]: from sklearn.metrics import classification_report  
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.49	0.65	68
1	0.48	1.00	0.65	32
accuracy			0.65	100
macro avg	0.74	0.74	0.65	100
weighted avg	0.83	0.65	0.65	100

```
[67]:
```