

```
In [1]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
%matplotlib inline
```

```
In [2]: df = pd.read_csv('Social_Network_Ads.csv')
df.head(10)
```

Out[2]:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19.0	19000.0	0
1	15810944	Male	35.0	20000.0	0
2	15668575	Female	26.0	43000.0	0
3	15603246	Female	27.0	57000.0	0
4	15804002	Male	19.0	76000.0	0
5	15728773	Male	27.0	58000.0	0
6	15598044	Female	27.0	84000.0	0
7	15694829	Female	32.0	150000.0	1
8	15600575	Male	25.0	33000.0	0
9	15727311	Female	35.0	65000.0	0

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User ID                400 non-null   int64
1   Gender                 400 non-null   object
2   Age                    400 non-null   float64
3   EstimatedSalary        400 non-null   float64
4   Purchased              400 non-null   int64
dtypes: float64(2), int64(2), object(1)
memory usage: 15.8+ KB
```

```
df.describe()
```

	User ID	Age	EstimatedSalary	Purchased
count	4.000000e+02	400.000000	400.000000	400.000000
mean	1.569154e+07	37.655000	69742.500000	0.357500
std	7.165832e+04	10.482877	34096.960282	0.479864
min	1.556669e+07	18.000000	15000.000000	0.000000
25%	1.562676e+07	29.750000	43000.000000	0.000000
50%	1.569434e+07	37.000000	70000.000000	0.000000
75%	1.575036e+07	46.000000	88000.000000	1.000000
max	1.581524e+07	60.000000	150000.000000	1.000000

```
x = df.iloc[:,[2,3]].values
y = df.iloc[:,4].values
```

X

```
array([[1.90e+01, 1.90e+04],  
       [3.50e+01, 2.00e+04],  
       [2.60e+01, 4.30e+04],  
       [2.70e+01, 5.70e+04],  
       [1.90e+01, 7.60e+04],  
       [2.70e+01, 5.80e+04],  
       [2.70e+01, 8.40e+04],  
       [3.20e+01, 1.50e+05],  
       [2.50e+01, 3.30e+04],  
       [3.50e+01, 6.50e+04],  
       [2.60e+01, 8.00e+04],  
       [2.60e+01, 5.20e+04],  
       [2.00e+01, 8.60e+04],  
       [3.20e+01, 1.80e+04],  
       [1.80e+01, 8.20e+04],  
       [2.90e+01, 8.00e+04],  
       [4.70e+01, 2.50e+04],  
       [4.50e+01, 2.60e+04],  
       [4.60e+01, 2.80e+04],  
       [4.80e+01, 2.80e+04]])
```

In [7]: y

Out[7]: array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1,
0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0,
1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0,
1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1,
0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1,
1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1,
0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0,
1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1,
0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1,
1, 1, 0, 1], dtype=int64)

Split data into train and test

In [9]: `from sklearn.model_selection import train_test_split`
`x_train , x_test , y_train , y_test = train_test_split(x,y,test_size = 0.25, random_`

Preprocessing

In [10]: `from sklearn.preprocessing import StandardScaler`
`sc = StandardScaler()`
`x_train = sc.fit_transform(x_train)`
`x_test = sc.transform(x_test)`

In [11]: x_train

```
[ 0.8787462 , -0.59677555],
[ 2.06713324, -1.17663843],
[ 1.07681071, -0.13288524],
[ 0.68068169,  1.78066227],
[-0.70576986,  0.56295021],
[ 0.77971394,  0.35999821],
[ 0.8787462 , -0.53878926],
[-1.20093113, -1.58254245],
[ 2.1661655 ,  0.93986109],
[-0.01254409,  1.22979253],
[ 0.18552042,  1.08482681],
[ 0.38358493, -0.48080297],
[-0.30964085, -0.30684411],
[ 0.97777845, -0.8287207 ],
[ 0.97777845,  1.8676417 ],
[-0.01254409,  1.25878567],
[-0.90383437,  2.27354572],
[-1.20093113, -1.58254245],
[ 2.1661655 , -0.79972756],
[-1.39899564, -1.46656987]
```

In [12]: `from sklearn.linear_model import LogisticRegression`
`classifier = LogisticRegression(random_state = 0)`
`classifier.fit(x_train,y_train)`

Out[12]: `LogisticRegression`
`LogisticRegression(random_state=0)`

Prediction

In [13]: `y_pred = classifier.predict(x_test)`

In [14]: `y_pred`

Out[14]: `array([0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,`
`0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,`
`1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,`
`0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1,`
`0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1], dtype=int64)`

Confusion Matrix

In [15]: `from sklearn.metrics import confusion_matrix,classification_report`
`cm = confusion_matrix(y_test , y_pred)`

In [16]: cm

Out[16]: array([[65, 3],
[8, 24]], dtype=int64)

In [17]: c1_report = classification_report(y_test,y_pred)

In [18]: c1_report

Out[18]: ' precision recall f1-score support\n\n 0 0.89
0.96 0.92 68\n 1 0.89 0.75 0.81 32\n\n accuracy 0.89 100\n macro avg 0.89
0.85 0.87 100\nweighted avg 0.89 0.89 0.89 100\n'

In [19]: tp, fn, fp, tn = confusion_matrix(y_test, y_pred, labels = [0,1]).reshape(-1)
print('Outcome values : \n', tp, fn, fp, tn)

Outcome values :
65 3 8 24

In [20]: accuracy_cm = (tp+tn)/(tp+fp+tn+fn)
precision_cm = tp/(tp+fp)
recall_cm = tp/(tp+fn)
f1_score = 2/((1/recall_cm)+(1/precision_cm))

In [21]: print("Accuracy : ",accuracy_cm)
print("Precision : ",precision_cm)
print("Recall : ",recall_cm)
print("F1-Score : ",f1_score)

Accuracy : 0.89
Precision : 0.8904109589041096
Recall : 0.9558823529411765
F1-Score : 0.9219858156028368

In []: